

BUSINESS ANALYTICS CLUB

Workshop Series 9.26



Machine Learning

Janet Ye
Shantanu Joshi

Set Ups

- Download the zip file from bit.ly/bacdata "Python ML" folder. In the zip file, you'll find:
 - Installation Guide
 - Python Installers
 - Sublime Text (Text editor) Installers
 - Windows Users:
 - Special Library Package File (IMPT: DO NOT edit the name of this file)

Learning Objective

1. Understand the fundamentals of a Neural Network and it's application through Python
2. Install and use basics of Python 2.7
3. Setup NumPy
4. Introduce Sigmoid Curves and Backpropagation

Terminology

Machine Learning: the study and construction of algorithms that can learn from and make predictions on data

Neural Network: statistical learning method inspired by biological neural networks

Backpropagation: how our Neural Network will learn...

Sigmoid Function: can map any number to be between 0,1

NumPy Setup

Numpy on Mac

- NumPy is a **library** in python that adds support for arrays or **data frames**
- Mac Users: Type **[Command+Space]** and type **[Terminal]**, hit enter to open a new terminal window
 - In the terminal window type **[pip install numpy]** and hit enter

Numpy on Windows

- Installation is tricky. Follow instructions carefully.
- Type [Windows+Q] and type [powershell], double click to open [Windows PowerShell]
- In PowerShell, type:
 `cd .\Downloads\python_mlworkshop\Windows\`
hit ENTER, then type:
 `pip install numpy`
and instead of hitting ENTER this time, hit the [TAB] key, this should autocomplete to:
 `numpy-1.9.3+mkl-cp27-none-win32.whl`
now hit ENTER
- Note: PowerShell will be referred to as the **Terminal** in the rest of the presentation.

Testing Numpy

- Open a new `Terminal` window (Windows: PowerShell)
- Type: `python`
- Once the Python prompt appears type: `import numpy as np` and hit ENTER, nothing should happen
- Type: `np` and hit ENTER the path of numpy should appear

Building our File

- Open Sublime Text and create a new file called `neural.py`
- Save to your DESKTOP
- The bottom right corner of sublime text should say `[Python]`
- Type: `print "Hello World"` and hit save
- Open `Terminal`, type `cd Desktop`. You should see:
Mac: `Desktop your_username$`
Windows: `C:\Users\your_username`

Running a Python File

Once you are in Desktop directory...

- Typing `[ls]` on mac or `[dir]` in windows will display `[neural.py]`


Running a file:

- Type: `python neural.py`
- You should see `Hello World` printed on your screen

Sigmoid Function

Sigmoid Function

- Maps any value to a value between 0 or 1
- Can convert numbers to probabilities
- $\frac{dS}{dx} = x * (1 - x)$



$$S(x) = \frac{1}{1 + e^{-x}}$$

Importing Numpy

In Sublime, edit `neural.py` delete the previous text and leave only:

```
import numpy as np
```

Syntax:

- Importing NumPy in this way allows us to use all the prebuilt functions in the NumPy `package`
- Setting numpy `as np` simply allows us to use `np` instead of typing out `numpy`

Input the Sigmoid Curve

In Sublime, edit `neural.py` and type →

```
# sigmoid function
def nonlin(x,deriv=False):
    if(deriv==True):
        return x*(1-x)
    return 1/(1+np.exp(-x))
```

Syntax:

- `def`: tells python we are defining a function called `nonlin` with parameters `x` and `deriv`
- `deriv` is an `optional variable`, and it is set to `False` by default. You will see in a few slides how to change its value.

Python Indentation

VERY IMPORTANT!

- Indents in python are considered syntax.
- 1 indent = 4 spaces
- Nested items require an extra indent than their parent items.
- Incorrectly placed or missing indents lead to most python syntax compile errors.

Input the Sigmoid Curve

- `return` tells our function to output something and terminate the function
- The next line is an `if-statement`, if `deriv` is set to `True` then the `nested` statement is returned
- If `deriv` is `False`, it skips the `if` statement and the last line is returned

```
# sigmoid function
def nonlin(x,deriv=False):
    if(deriv==True):
        return x*(1-x)
    return 1/(1+np.exp(-x))
```


Note: NumPy Usage

- Notice the `np.exp(-x)`
- As previously defined we set numpy as np so this function actually reads as:
`numpy.exp(something)`
- The `numpy.exp` function is a function that returns the value of `[e]` (~2.718...) raised to the `[something]` power

```
# sigmoid function
def nonlin(x,deriv=False):
    if(deriv==True):
        return x*(1-x)
    return 1/(1+np.exp(-x))
```

Sigmoid Defined

- In the function `nonlin` we've defined the behavior of our `sigmoid curve`
- When `[deriv=false]` the function tells us $\frac{1}{1 + e^{-x}}$
- When `[deriv=true]` the function outputs the derivative at a given `[x]` which is: $1 * (1 - x)$

Initializing the Dataset

Our Data Set

Inputs:	(1)	(2)	(3)	Output
	0	0	1	0
	1	1	1	1
	1	0	1	1
	0	1	1	0

- We are trying to predict the Output column via the three Input columns
- Measuring statistics ex) Leftmost column is perfectly correlated with output will help solve the problem

Initializing the Inputs

- Type the following after the `nonlin` definition:

```
# input dataset
X = np.array([ [0,0,1],
               [0,1,1],
               [1,0,1],
               [1,1,1] ])
```

Syntax:

- We are setting `X` to be a matrix representing our `inputs`
- `np.array` is the NumPy function for multidimensional arrays
- Can also be written in one line (no need of special spacing)

```
X = np.array([ [0,0,1], [0,1,1], [1,0,1], [1,1,1] ])
```

Initializing the Output

- Type the following after the `input` definition:

```
y = np.array([[0,0,1,1]]).T
```

Syntax:

- We are setting `y` (note: lowercase) to be a vector representing our `output`
- The `.T` at the end of the array definition is the transpose function in numpy
- Why? Because It's easier to write `.T` then `np.array([[0],[0],[1],[1]])`

Array Transpose...

$$y = [0 \quad 0 \quad 1 \quad 1]$$

$$y^T = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

Initializing Random Seed

- Type the following after the `output` definition:

```
np.random.seed(1)
```

Syntax:

- This refers to the seed number for the `random function` in the `np.random module`
- Random numbers come from normal distributions, the seed allows each random number to come from the SAME normal distribution
- That way we train the model using the same random numbers. This allows us to see the true effects of our neural network.

Initializing Weights

Initializing Weights

- Type the following after the `random seed` function:

```
# initialize weights randomly with mean 0  
syn0 = 2*np.random.random((3,1)) - 1
```

Syntax:

- Initial weights will be set to `syn0`, the first layer synapse
- `np.random.random((r,c))` returns an `r` by `c` matrix (in this case 3x1) of random numbers in the range `[0, 1)`, based off of the seed
- `2 * np.random.random - 1` manipulates the interval and returns random numbers in the range `[-1, 1)`

Understand Neural Networks


Human Neurons



blue neurons / neu / January 2010.

Neural Network

- Neural network is a machine learning algorithm that tries to **mimic** the human brain.
- Widely used in the 80s and early 90s, but popularity diminished because it requires intense computational power.
- Recent resurgence: state-of-the-art technique for many applications



Training Session

```
for iter in xrange(10000):  
  
    # forward propagation  
    l0 = X  
    l1 = nonlin(np.dot(l0,syn0))  
  
    # how much did we miss?  
    l1_error = y - l1  
  
    # multiply how much we missed by the  
    # slope of the sigmoid at the values in l1  
    l1_delta = l1_error * nonlin(l1,deriv=True)  
  
    # update weights  
    syn0 += np.dot(l0.T,l1_delta)
```

Breaking it down line by line...

Training Session

- We begin a loop by writing: `for iter in xrange(10000):`
- (Don't miss the colon at the end!)
- `xrange(10000)` means we are looping `iter`, our placeholder variable, from 0 to 9999. You can name it anything you'd like, usually people use something like `i`, or `iter`
- The loop first sets `iter` to 0, computes everything that is indented under the for loop.
- Then set `iter` to 1, compute...set to 2, compute...etc. The loop stops after `iter` completes the 10000th round. The value of `iter` at completion is 9999.

Training Session

- Forward propagation: going forward with the layers
- Sublime text should automatically begin indented block within the for loop
- `l0 = X` here we set the input layer of the neural network to X
- `l1 = nonlin(np.dot(l0, syn0))` computes the hidden layer:
 - `np.dot(l0, syn0)` is the dot product calculation for matrices
 - Then we pass the dot product to the `nonlin` function we wrote
- How did we do? `l1_error = y - l1` compares our prediction (what sigmoid function outputs) with the actual y

Training Session

```
l1_delta = l1_error * nonlin(l1,deriv=True)
```

- Recall in Calc 1, Euler's method says $dy = dx * \text{slope}$
- `l1_error` is like your `dx`
- `nonlin(l1, True)` calls the `nonlin` function. It sees the `deriv=True`, and enters into the if statement

```
# sigmoid function
def nonlin(x,deriv=False):
    if(deriv==True):
        return x*(1-x)
    return 1/(1+np.exp(-x))
```

- `l1_delta` is like your `dy`

Training Session

```
syn0 += np.dot(l0.T, l1_delta)
```

- Computer science syntax: $i += 1$ is the shorthand for $i = i + 1$ (Similarly, $i -= 1$ is the same as $i = i - 1$)
- We now adjust `syn0`, the first layer synapse, by adding an adjustment term – the dot product of `l0` and `l1_delta`
- This step is called `backward propogation`
- The loop continues...`iter` is incremented by 1

- So far, you should have something like this.
- Take a second to modify your code if need be.
- Pay attention to indentation, commas, and colons

```
import numpy as np

# sigmoid function
def nonlin(x,deriv=False):
    if(deriv==True):
        return x*(1-x)
    return 1/(1+np.exp(-x))

# input dataset
X = np.array([ [0,0,1],
               [0,1,1],
               [1,0,1],
               [1,1,1] ])

# output dataset
y = np.array([[0,0,1,1]]).T

# seed random numbers to make calculation
# deterministic (just a good practice)
np.random.seed(1)

# initialize weights randomly with mean 0
syn0 = 2*np.random.random((3,1)) - 1

for iter in xrange(10000):

    # forward propagation
    l0 = X
    l1 = nonlin(np.dot(l0,syn0))

    # how much did we miss?
    l1_error = y - l1

    # multiply how much we missed by the
    # slope of the sigmoid at the values in l1
    l1_delta = l1_error * nonlin(l1,deriv=True)

    # update weights
    syn0 += np.dot(l0.T,l1_delta)

print "Output After Training:"
print l1
```

Running Code

- In Terminal, (Desktop should be your current directory), type `python neural.py`
- You should get something like...results are quite good

```
MacBook:ml shantanu$ python neural.py
Output After Training:
[[ 0.00966449]
 [ 0.00786506]
 [ 0.99358898]
 [ 0.99211957]]
```

Real Life Example: Self-Driving Car

Useful Resources

- Free Stanford Machine Learning on Coursera
 - Link: <https://www.coursera.org/learn/machine-learning>
 - Blog with detailed write-ups
<http://www.holehouse.org/mlclass/index.html>
- Reddit: <https://www.reddit.com/r/MachineLearning/wiki/index>
- Machine Learning Data Sets: <http://archive.ics.uci.edu/ml/>
- Machine Learning Visualization: <http://www.r2d3.us/visual-intro-to-machine-learning-part-1/>
- Dope Blog: <http://www.wzchen.com>

Acknowledgement

Andrew Trask

Product Manager @
Digital Reasoning

Appendix – Command Line Navigation

Navigation Tips:

- Type `[cd]` to change directories
- `[cd ..]` to move up a directory
- Type `[ls]` on mac or `[dir]` in windows to see contents of your current directory
- Typing `tab` will autocomplete with items from your current directory
- Typing `[pwd]` on a mac or `[cd]` in windows will list the current directory path